
Big Data Processing

M. Tamer Özsu
Patrick Valduriez

Outline

- Big Data Processing
 - ❑ Distributed storage systems
 - ❑ Processing platforms
 - ❑ Stream data management
 - ❑ Graph analytics
 - ❑ Data lake

Four Vs

■ Volume

- ▣ Increasing data size: petabytes (10^{15}) to zettabytes (10^{21})

■ Variety

- ▣ Multimodal data: structured, images, text, audio, video
- ▣ 90% of currently generated data unstructured

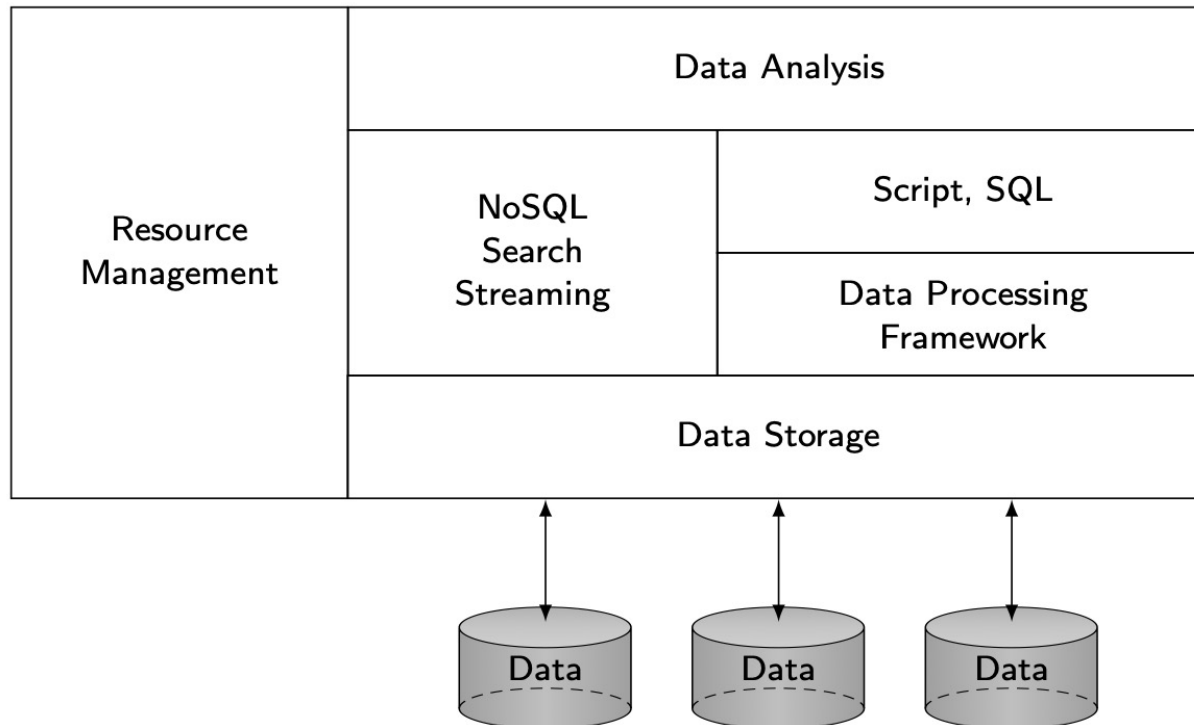
■ Velocity

- ▣ Streaming data at high speed
- ▣ Real-time processing

■ Veracity

- ▣ Data quality

Big Data Software Stack



Outline

- Big Data Processing
 - Distributed storage systems
 - Processing platforms
 - Stream data management
 - Graph analytics

Distributed Storage System

Storing and managing data across the nodes of a shared-nothing cluster

■ Object-based

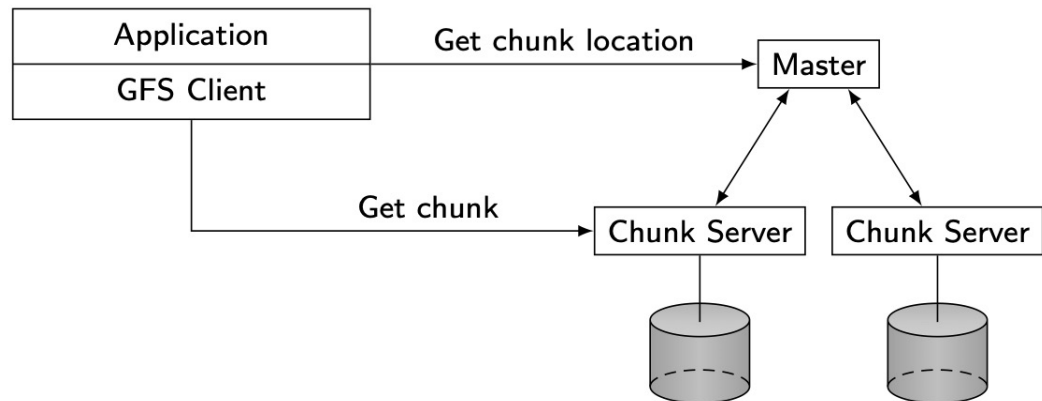
- ❑ Object = $\langle \text{oid}, \text{data}, \text{metadata} \rangle$
- ❑ Metadata can be different for different object
- ❑ Easy to move
- ❑ Flat object space \rightarrow billions/trillions of objects
- ❑ Easily accessed through REST-based API (get/put)
- ❑ Good for high number of small objects (photos, mail attachments)

■ File-based

- ❑ Data in files of fixed- or variable-length records
- ❑ Metadata-per-file stored separately from file
- ❑ For large data, a file needs to be partitioned and distributed

Google File System (GFS)

- Targets shared-nothing clusters of thousands of machines
- Targets applications with characteristics:
 - ❑ Very large files (several gigabytes)
 - ❑ Mostly read and append workloads
 - ❑ High throughput more important than low latency
- Interface: create, open, read, write, close, delete, snapshot, record append



Outline

- Big Data Processing
 - Distributed storage systems
 - Processing platforms
 - Stream data management
 - Graph analytics

Big Data Processing Platforms

- Applications that do not need full DBMS functionality
 - Data analysis of very large data sets
 - Highly dynamic, irregular, schemaless, ...
- “Embarrassingly parallel problems”
- MapReduce/Spark
- Advantages
 - Flexibility
 - Scalability
 - Efficiency
 - Fault-tolerance
- Disadvantage
 - Reduced functionality
 - Increased programming effort

MapReduce Basics

- Simple programming model
 - Data structured as (key, value) pairs
 - E.g. (doc-id, content); (word, count)
 - Functional programming style with two functions
 - $\text{map}(k1, v1) \rightarrow \text{list}(k2, v2)$
 - $\text{reduce}(k2, \text{list}(v2)) \rightarrow \text{list}(v3)$
- Implemented on a distributed file system (e.g. Google File System) on very large clusters

map Function

- User-defined function
 - ❑ Processes input (key, value) pairs
 - ❑ Produces a set of **intermediate** (key, value) pairs
 - ❑ Executes on multiple machines (called **mapper**)
- map function I/O
 - ❑ **Input:** read a **chunk** from distributed file system (DFS)
 - ❑ **Output:** Write to intermediate file on local disk
- MapReduce library
 - ❑ Execute map function
 - ❑ Groups together all intermediate values with same key
 - ❑ Passes these lists to reduce function
- Effect of map function
 - ❑ Processes and partitions input data
 - ❑ Builds a distributed map (transparent to user)
 - ❑ Similar to “group by” operator in SQL

reduce Function

- User-defined function
 - ❑ Accepts one intermediate key and a set of values for that key (i.e. a list)
 - ❑ Merges these values together to form a (possibly) smaller set
 - ❑ Computes the reduce function generating, typically, zero or one output per invocation
 - ❑ Executes on multiple machines (called **reducer**)
- reduce function I/O
 - ❑ **Input:** read from intermediate files using remote reads on local files of corresponding mappers
 - ❑ **Output:** Write result back to DFS
- Effect of map function
 - ❑ Similar to aggregation function in SQL

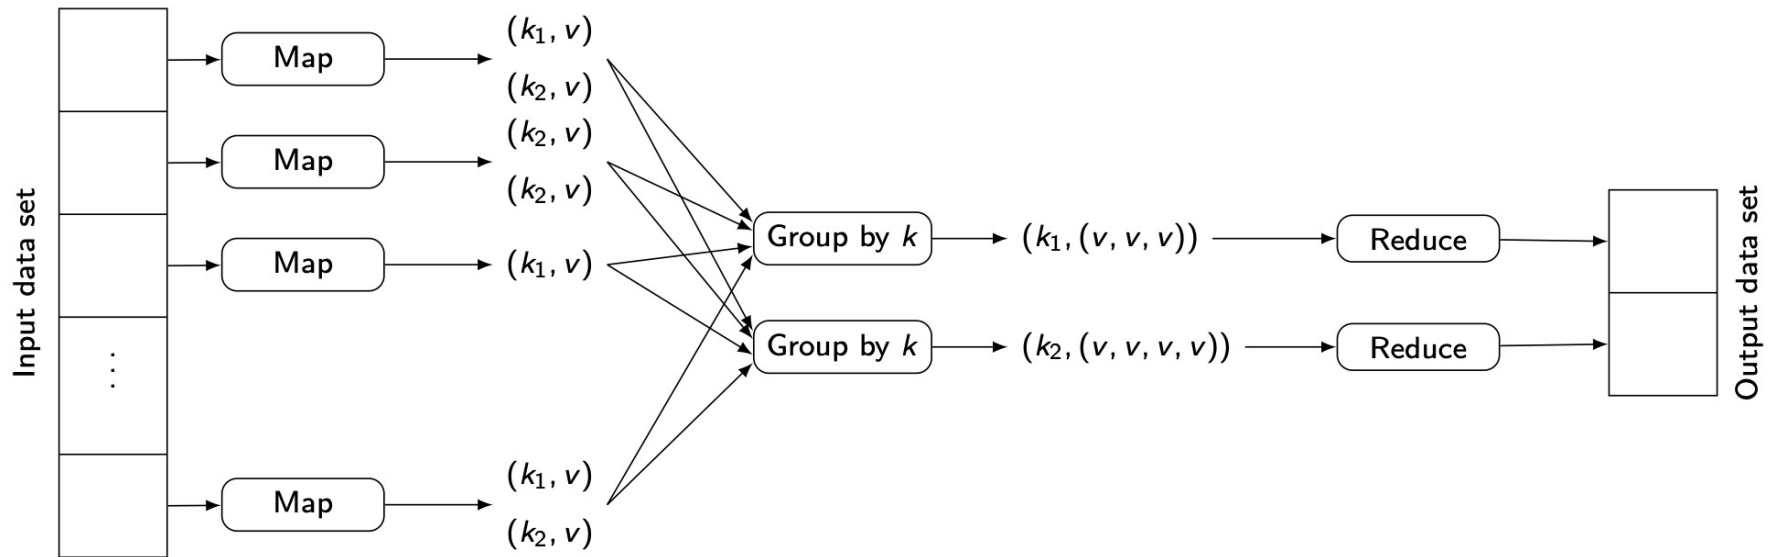
Example

Consider EMP (ENO, ENAME, TITLE, CITY)

```
SELECT      CITY, COUNT (*)  
FROM        EMP  
WHERE        ENAME LIKE "%Smith"  
GROUP BY    CITY
```

```
map (Input: (TID,EMP), Output: (CITY, 1)  
    if EMP.ENAME like ``\%Smith'' return (CITY, 1)  
reduce (Input: (CITY, list(1)), Output: (CITY,  
SUM(list)))  
    return (CITY, SUM(1))
```

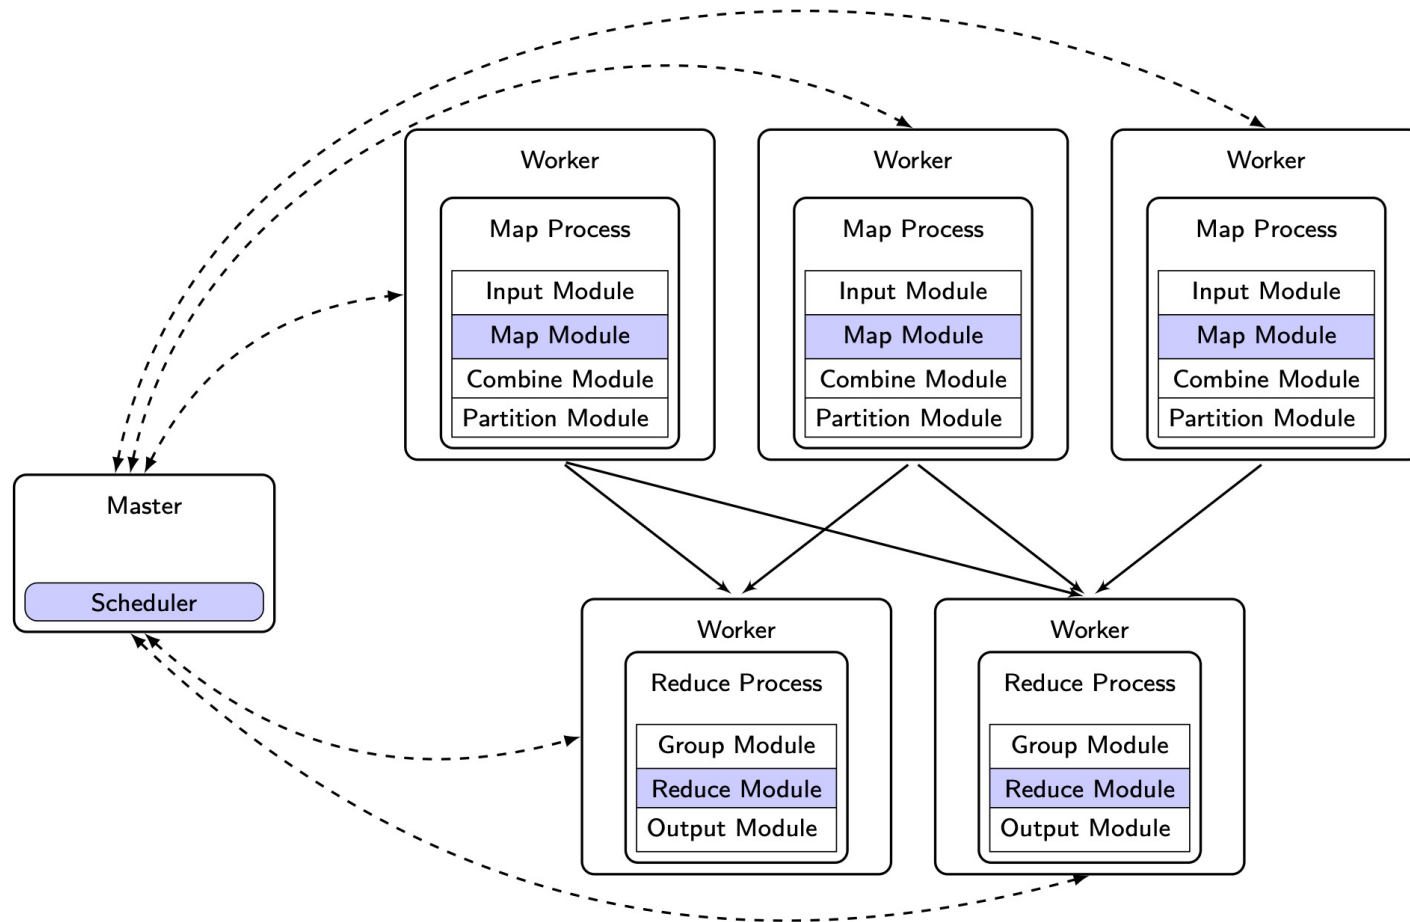
MapReduce Processing



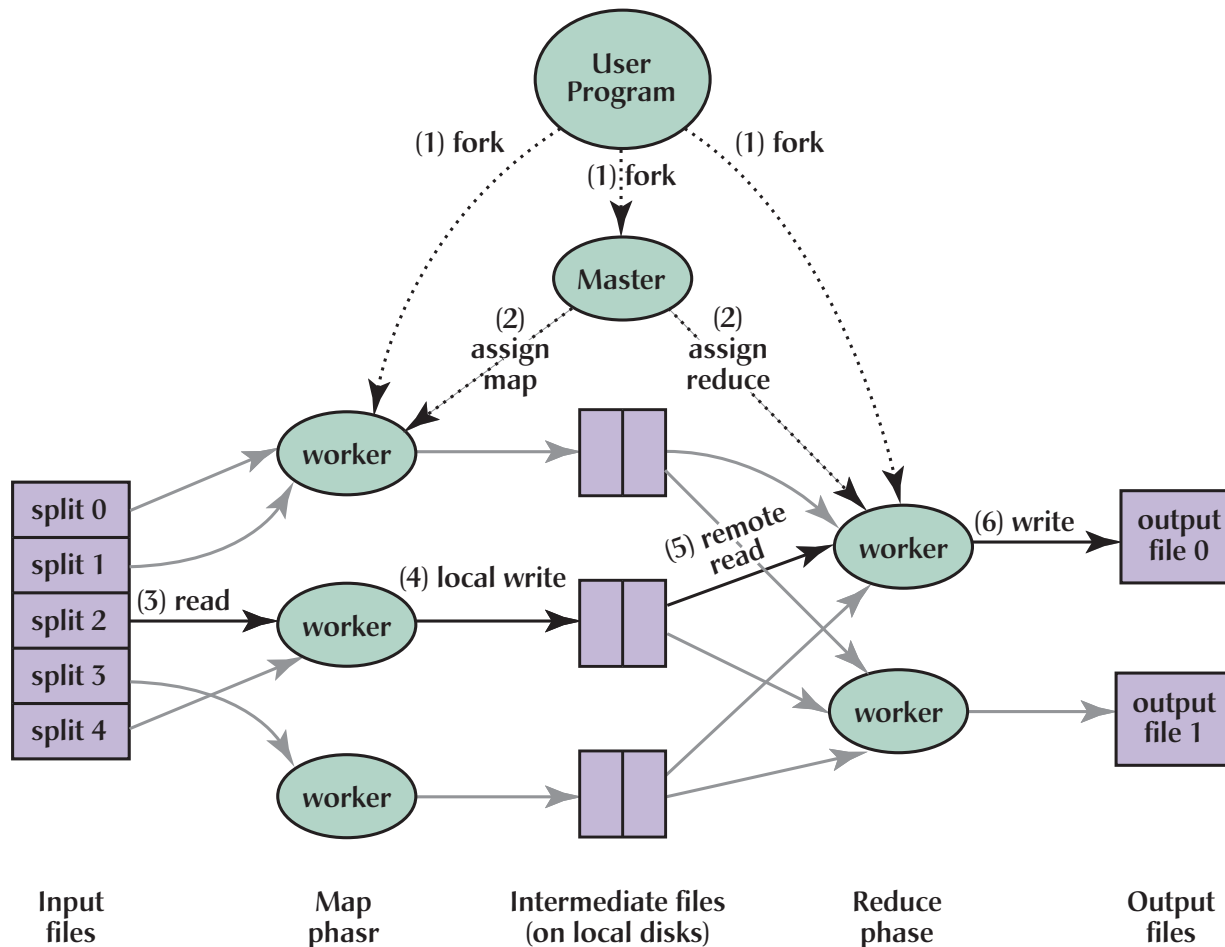
Hadoop Stack

Yarn	Third party analysis tools R (statistics), Mahout (machine learning), ...	
	Hbase	Hive & HiveQL
		Hadoop (MapReduce engine)
	Hadoop Distributed File System (HDFS)	

Master-Worker Architecture

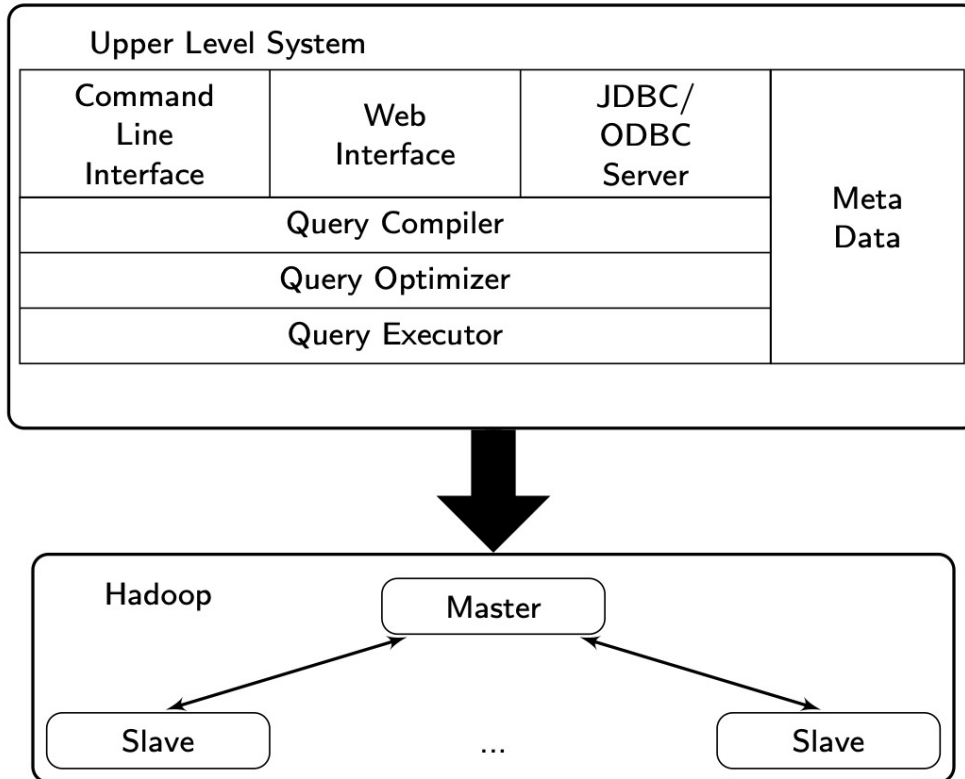


Execution Flow



From: J. Dean and S. Ghemawat. MapReduce: Simplified data processing on large clusters, *Comm. ACM*, 51(1), 2008.

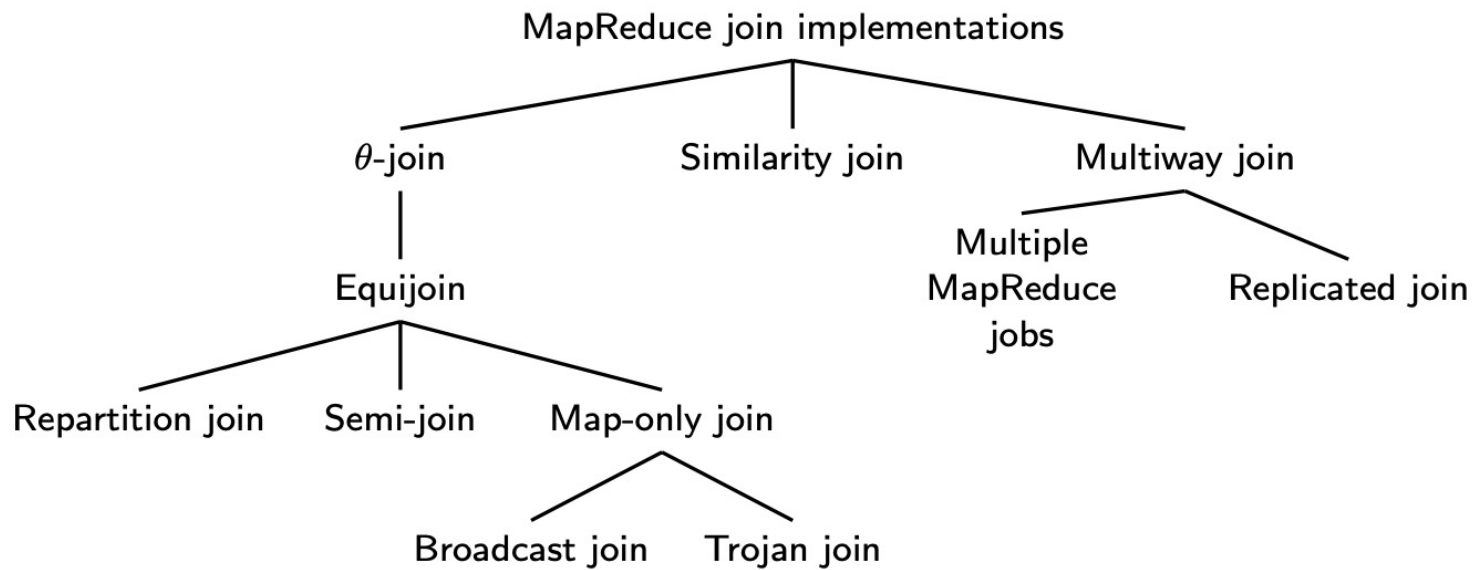
High-Level MapReduce Languages



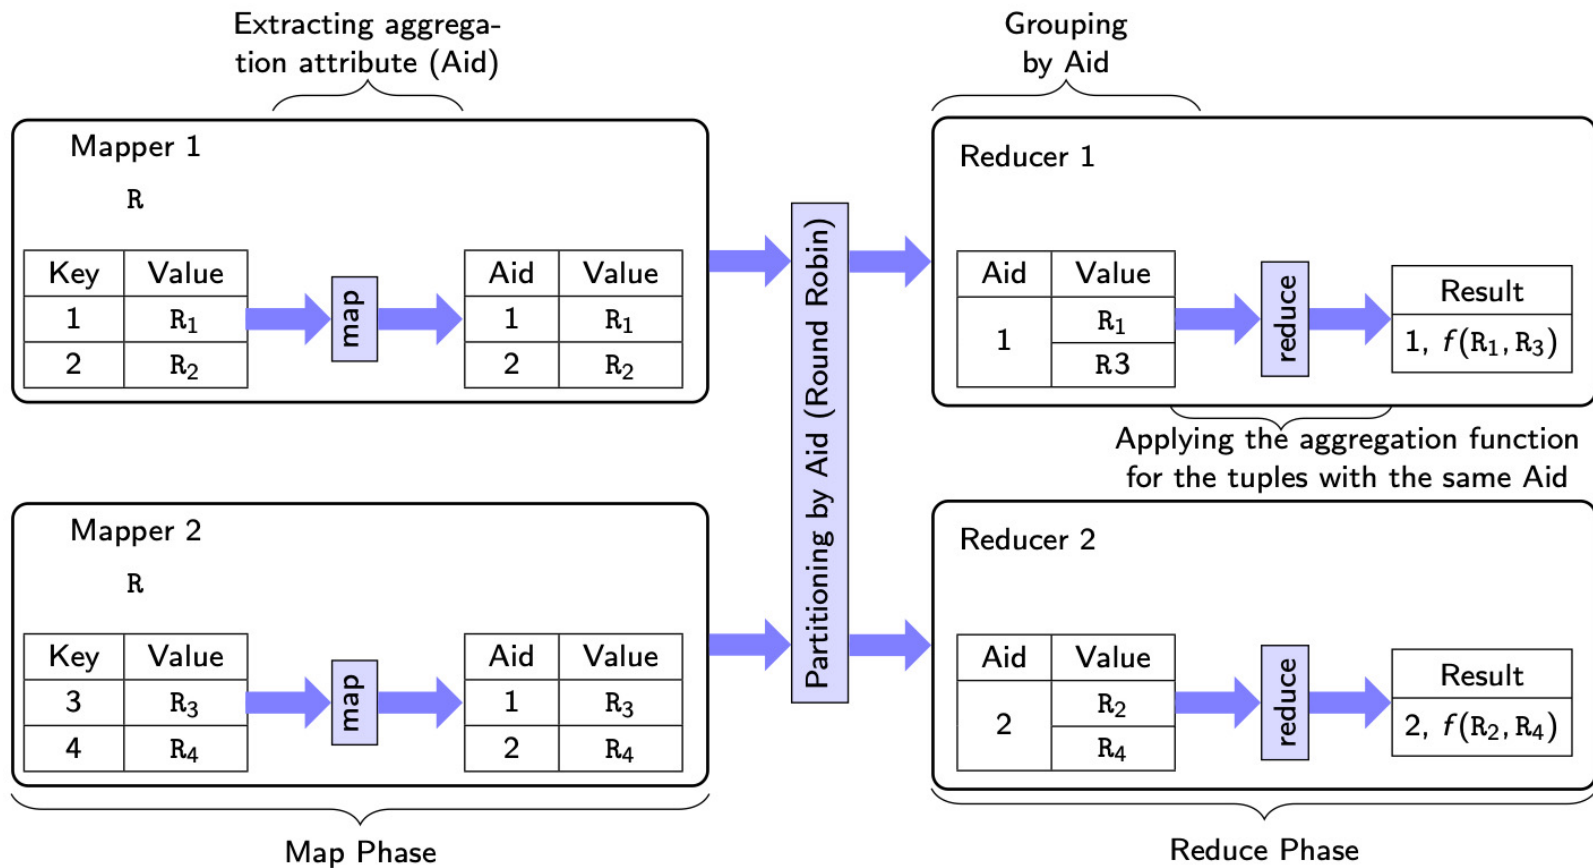
- Declarative
 - ❑ HiveQL
 - ❑ Tenzing
 - ❑ JAQL
- Data flow
 - ❑ Pig Latin
- Procedural
 - ❑ Sawzall
- Java Library
 - ❑ FlumeJava

MapReduce Implementations of DB Ops

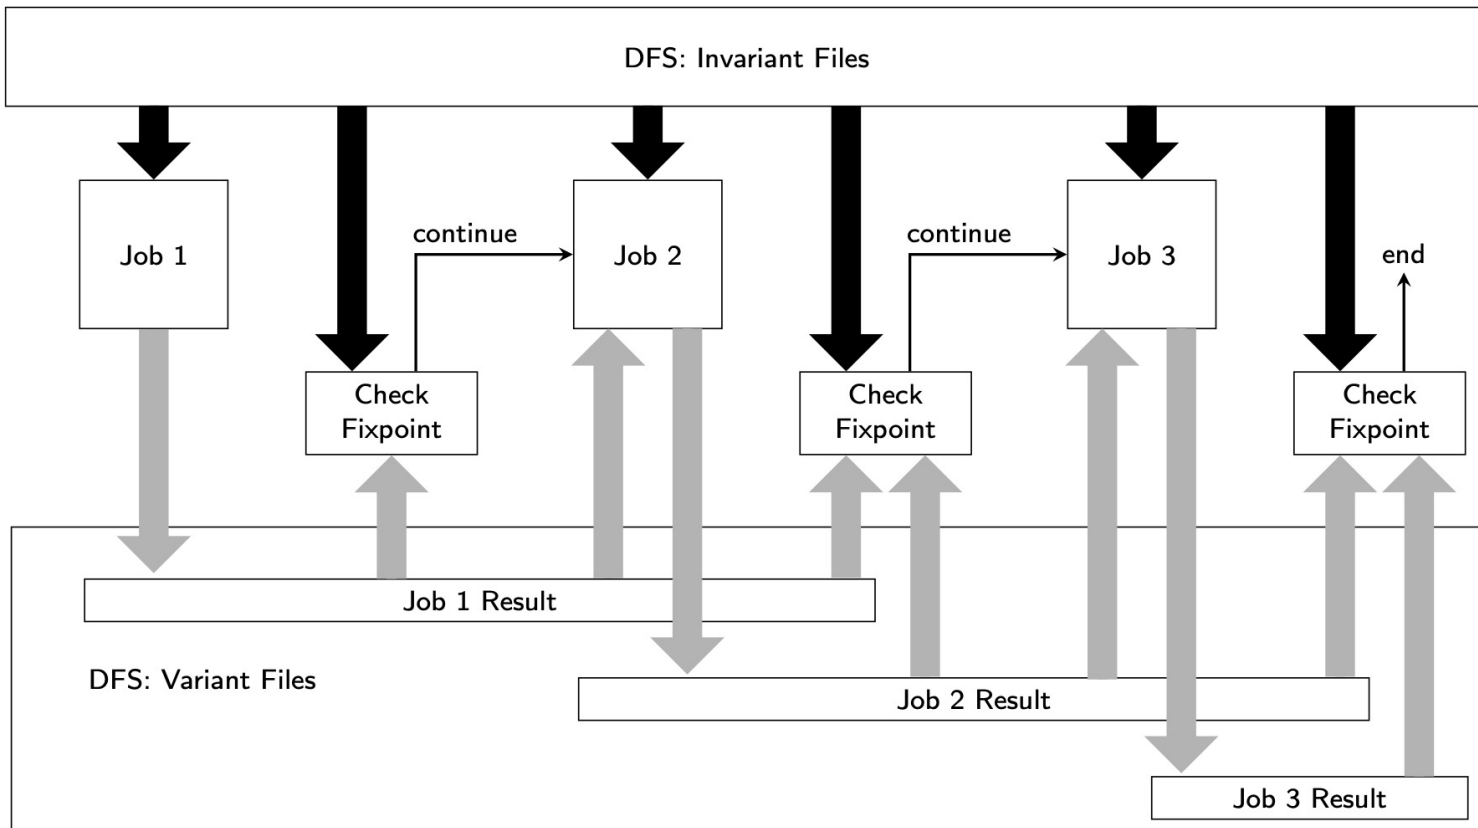
- Select and Project can be easily implemented in the map function
- Aggregation is not difficult (see next slide)
- Join requires more work



Aggregation



MapReduce Iterative Computation



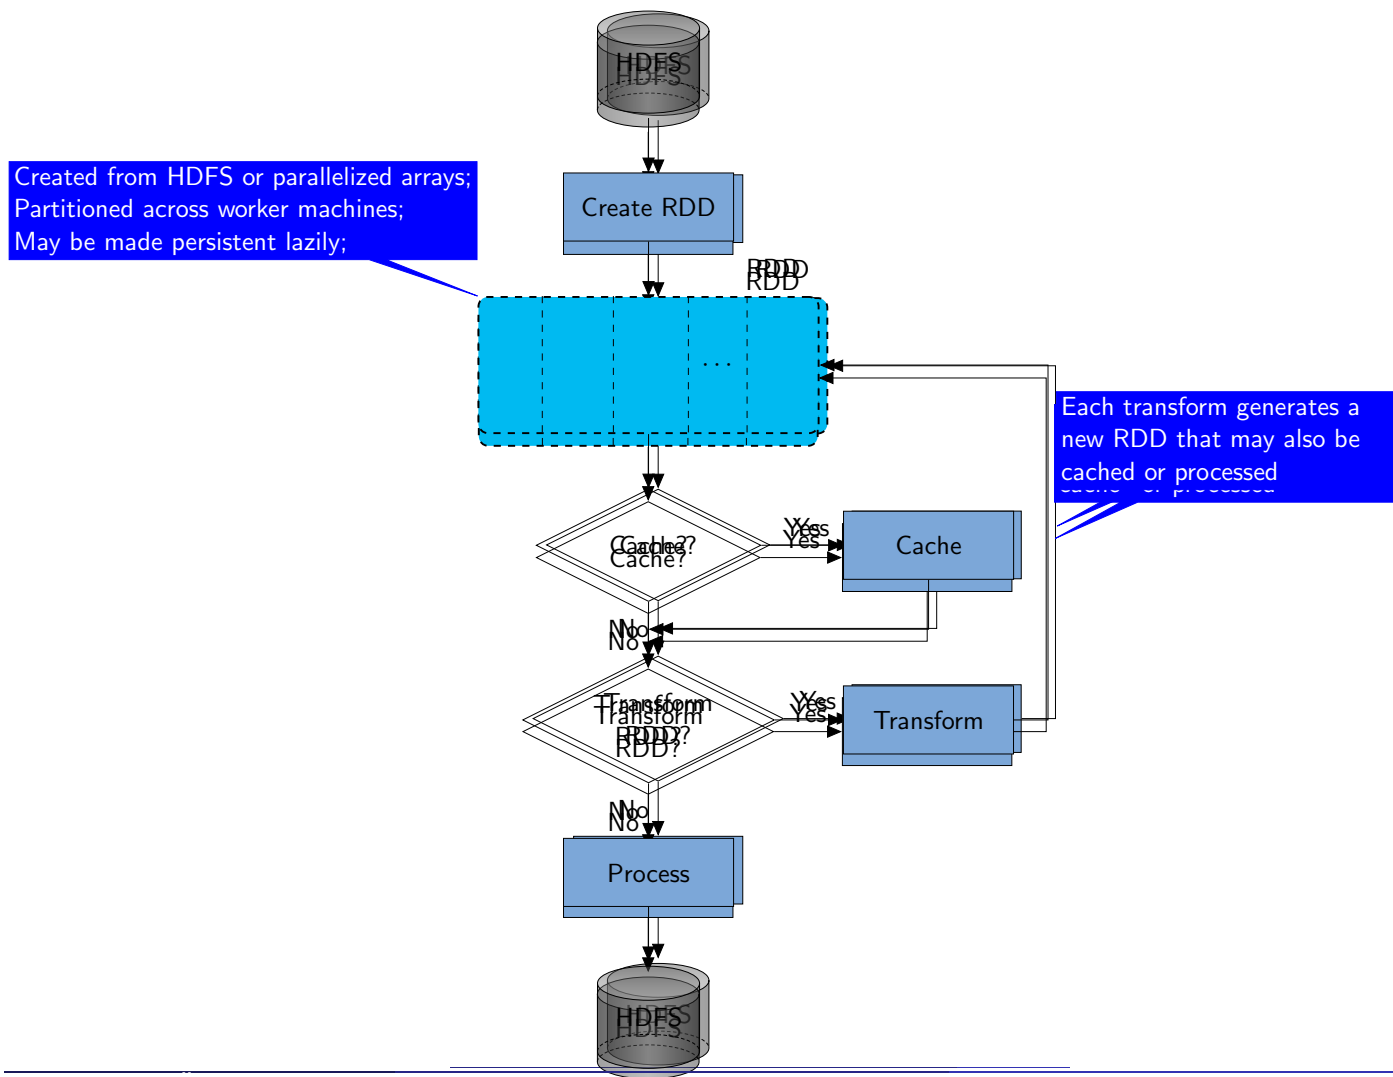
Problems with Iteration

- MapReduce workflow model is acyclic
 - ▣ Iteration: Intermediate results have to be written to HDFS after each iteration and read again
- At each iteration, no guarantee that the same job is assigned to the same compute node
 - ▣ Invariant files cannot be locally cached
- Check for fixpoint
 - ▣ At the end of each iteration, another job is needed

Spark

- Addresses MapReduce shortcomings
- Data sharing abstraction: Resilient Distributed Dataset (RDD)
 - 1) Cache working set (i.e. RDDs) so no writing-to/reading-from HDFS
 - 2) Assign partitions to the same machine across iterations
 - 3) Maintain lineage for fault-tolerance



Spark Program Flow



Outline

- Big Data Processing
 - Distributed storage systems
 - Processing platforms
 - Stream data management
 - Graph analytics

Traditional DBMS vs Streaming



■ Other differences

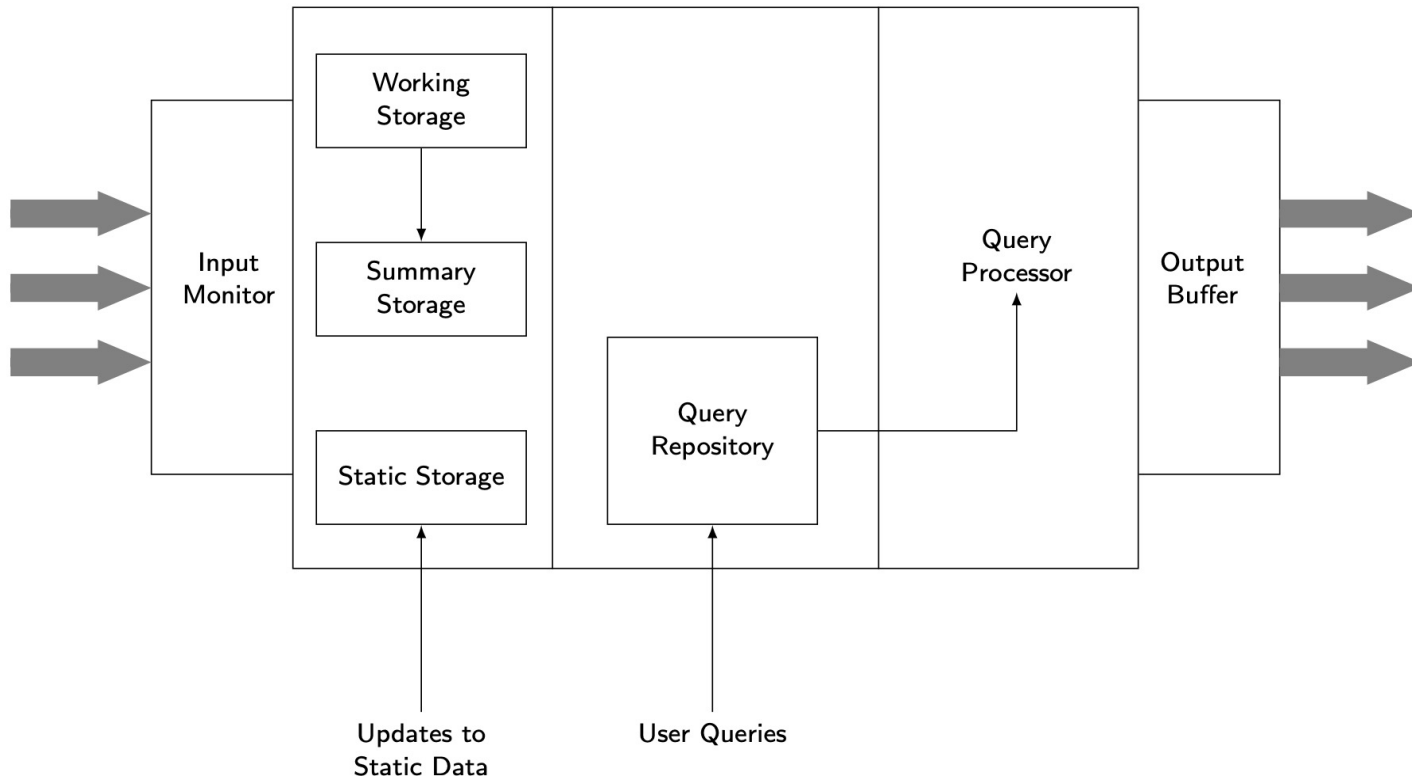
- ❑ Push-based (data-driven)
- ❑ Persistent queries

- ❑ Unbounded stream
- ❑ System conditions may not be stable

History

- **Data Stream Management System (DSMS)**
 - ❑ Typical DBMS functionality, primarily query language
 - ❑ Earlier systems: STREAM, Gigascope, TelegraphCQ, Aurora, Borealis
 - ❑ Mostly single machine (except Borealis)
- **Data Stream Processing System (DSPS)**
 - ❑ Do not embody DBMS functionality
 - ❑ Later systems: Apache Storm, Heron, Spark Streaming, Flink, MillWheel, TimeStream
 - ❑ Almost all are distributed/parallel systems
- Use **Data Stream System (DSS)** when the distinction is not important

DSMS Architecture



Stream Data Model

- Standard def: An append-only sequence of timestamped items that arrive in some order
- Relaxations
 - Revision tuples
 - Sequence of events that are reported continually (publish/subscribe systems)
 - Sequence of sets of elements (bursty arrivals)
- Typical arrival:

$\langle \text{timestamp}, \text{payload} \rangle$

 - Payload changes based on system
 - Relational: tuple
 - Graph: edge
 - ...

Processing Models

■ Continuous

- ❑ Each new arrival is processed as soon as it arrives in the system.
- ❑ Examples: Apache Storm, Heron

■ Windowed

- ❑ Arrivals are batched in windows and executed as a batch.
- ❑ For user, recently arrived data may be more interesting and useful.
- ❑ Examples: Aurora, STREAM, Spark Streaming

Stream Query Models

- Queries are typically persistent
- They may be **monotonic** or **non-monotonic**
- Monotonic: result set always grows
 - Results can be updated incrementally
 - Answer is continuous, append-only stream of results
 - Results may be removed from the answer only by explicit deletions (if allowed)
- Non-monotonic: some answers in the result set become invalid with new arrivals
 - Recomputation may be necessary

Stream Query Languages

■ Declarative

- ❑ SQL-like syntax, stream-specific semantics
- ❑ Examples: CQL, GSQL, StreaQuel

■ Procedural

- ❑ Construct queries by defining an acyclic graph of operators
- ❑ Example: Aurora

■ Windowed languages

- ❑ `size`: window length
- ❑ `slide`: how frequently the window moves
- ❑ E.g.: `size=10min, slide=5sec`

■ Monotonic vs non-monotonic

Outline

■ Big Data Processing

- ❑ Distributed storage systems
- ❑ Processing platforms
- ❑ Stream data management
- ❑ ~~Graph analytics~~

Data Lake

- Collection of raw data in native format
 - ▣ Each element has a unique identifier and metadata
 - ▣ For each business question, you can find the relevant data set to analyze it
- Originally based on Hadoop
 - ▣ Enterprise Hadoop

Advantages of a Data Lake

■ Schema on read

- ❑ Write the data as they are, read them according to a diagram (e.g. code of the Map function)
- ❑ More flexibility, multiple views of the same data

■ Multi-workload data processing

- ❑ Different types of processing on the same data
- ❑ Interactive, batch, real time

■ Cost-effective data architecture

- ❑ Excellent cost/performance and ROI ratio with SN cluster and open source technologies

Principles of a Data Lake

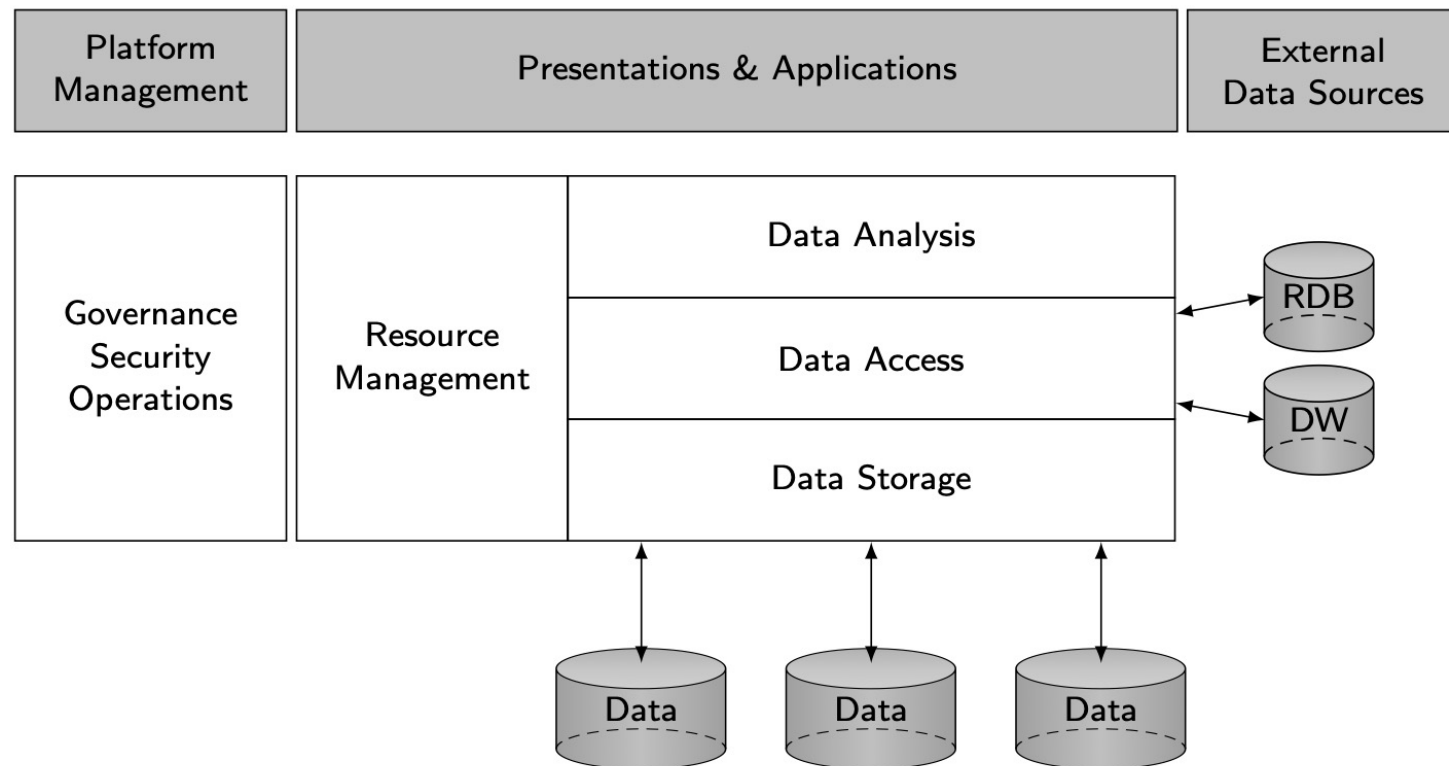
- Collect all useful data
 - Raw data, transformed data
- Dive from anywhere
 - Users from different business units can explore and enrich the data
- Flexible access
 - Different access paths to shared infrastructure
 - Batch, interactive (OLAP and BI), real-time, search,.....

Main Functions

- Data management, to store and process large amounts of data
- Data access: interactive, batch, real time, streaming
- Governance: load data easily, and manage it according to a policy implemented by the *data steward*
- Security: authentication, access control, data protection
- Platform management: provision, monitoring and scheduling of tasks (in a cluster)

Data Lake Architecture

A collection of multi-modal data stored in their raw formats



Data Lake vs Data Warehouse

Data Lake

- Shorter development process
- Schema-on-read
- Multiworkload processing
- Cost-effective architecture

Data Warehouse

- Long development process
- Schema-on-write
- OLAP workloads
- Complex development with ETL